

PySpark & Spark

Questions and answers from many sources — de-duplicated and organised by section

Questions	2
Sections	1
AI-written answers	0 (marked  — the source had the question only)
Last updated	2026-09-30

Sources merged into this file

- sql_spark_notes.pdf — 2

Contents

Performance & Tuning (2)	3
------------------------------------	---

Performance & Tuning (2)

Q1. What is a broadcast join in Spark?

The small table is copied to every executor so the large table does not need to be shuffled. Use it when one side comfortably fits in memory.

Source: [sql_spark_notes.pdf](#)

Q2. Why is data skew a problem in Spark?

A few keys hold most rows, so one task runs far longer than the rest. Fix with AQE skew join, salting, or broadcasting the smaller side.

Source: [sql_spark_notes.pdf](#)
